

dog_app

February 13, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/"))
        dog_files = np.array(glob("/data/dog_images/*/"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[9846])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

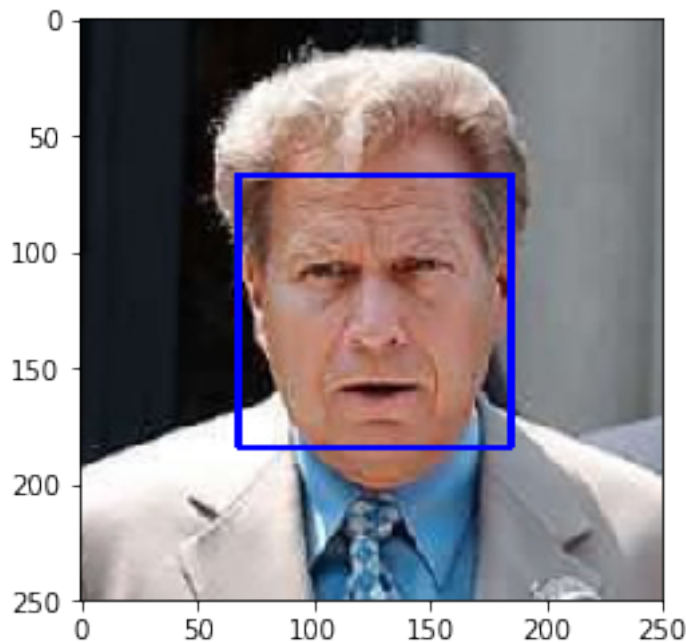
# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer:

98% Percentage of human face have been detected in `human_files`
17% Percentage of the Human faces have been detected in `Dog_files`

```
In [4]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
human_face_percent = []
dog_found = []

# Human Faces!!
for faces in human_files_short:
    human_face_percent.append(face_detector(faces))
    percent = np.sum(human_face_percent)
print ('{}% Percentage of human face have been detected in human_files'.format(percent))

# For dogs_files
for dogs in dog_files_short:
    dog_found.append(face_detector(dogs))
    d_percent = np.sum(dog_found)
```

```
print ('{}% Percentage of the Human faces have been detected in Dog_files '.format(d_pe
```

98% Percentage of human face have been detected in human_files

17% Percentage of the Human faces have been detected in Dog_files

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on human_files_short and dog_files_short.

```
In [ ]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [5]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()
        # Print the model
        print (VGG16)
```

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth
100%|| 553433881/553433881 [00:30<00:00, 17887801.99it/s]

```

VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): ReLU(inplace)
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (27): ReLU(inplace)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace)
    (2): Dropout(p=0.5)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace)
    (5): Dropout(p=0.5)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
)

```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [6]: from PIL import Image
import torchvision.transforms as transforms
from torchvision import datasets
def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """

    ## TODO: Complete the function.

    ## Load and pre-process an image from the given img_path

    transform = transforms.Compose([transforms.Resize(224),
                                    transforms.CenterCrop(224),
                                    transforms.ToTensor()
                                    ])

    def image_load(img_path):
        """ Load Image __ Return Tensor """
        #Open the Image "Using Image.open"
        image = Image.open(img_path)
        #Transform the image "Define above"
        image = transform(image).float()
        image = image.unsqueeze(0)
        if torch.cuda.is_available():
            image = image.cuda()
        return image

    # Move the images to cuda

    ## Return the *index* of the predicted class for that image
```

```

out = VGG16(image_load(img_path))
pred = torch.argmax(out)

''' Hello Mr. Reviewer can you teach me the difference between torch.argmax and torch.max

# _, preds = torch.max(out, 1)
#print (_)

return pred
# predicted class index

#VGG16_predict('/data/dog_images/train/001.Affenpinscher/Affenpinscher_00001.jpg')
#

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [7]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    pred = VGG16_predict(img_path)

    #Compare if dog the return true else false

    def compare(output):
        if output >= 151 and output < 269:
            return True
        else:
            return False

    return compare(pred) # true/false
#

```

```

In [8]: # Tests for Dogs
print (dog_detector('/data/dog_images/train/005.Alaskan_malamute/Alaskan_malamute_00300.jpg'))
print (dog_detector('/data/dog_images/train/084.Icelandic_sheepdog/Icelandic_sheepdog_005.jpg'))
print (dog_detector('/data/dog_images/train/004.Akita/Akita_00269.jpg'))
print (dog_detector('/data/dog_images/train/026.Black_russian_terrier/Black_russian_terrier_00269.jpg'))
print (dog_detector('/data/dog_images/train/003.Airedale_terrier/Airedale_terrier_00154.jpg'))

```


True
True
True
True
True

```
In [9]: ### Test for Humans
        print(dog_detector('/data/lfw/Adam_Sandler/Adam_Sandler_0002.jpg')) #LOL :D
        print (dog_detector('/data/lfw/Anil_Ramsook/Anil_Ramsook_0001.jpg')) # No idea :P
        print (dog_detector('/data/lfw/Bernadette_Peters/Bernadette_Peters_0001.jpg'))
```

False
False
False

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your dog_detector function.

- What percentage of the images in human_files_short have a detected dog?
- What percentage of the images in dog_files_short have a detected dog?

Answer:

0% percentage of the images in human_files_short have a detected dog

95% percentage of the images in dog_files_short have a detected dog

```
In [10]: ### TODO: Test the performance of the dog_detector function!! Let's set the Bar high!!
```

```
    # For Humans!! Let's Do this!!
    percent_h = []
    for human in human_files_short:
        percent_h.append(dog_detector(human))
    print ('{}% percentage of the images in human_files_short have a detected dog'.format(np.sum(percent_h)/len(percent_h)*100))
```

```
    ### on the images in human_files_short and dog_files_short.
```

```
    # For Dogs!! Who let the dogs out!!
```

```
    percent_d = []
    for dog_ in dog_files_short:
        percent_d.append(dog_detector(dog_))
    print('{}% percentage of the images in dog_files_short have a detected dog'.format(np.sum(percent_d)/len(percent_d)*100))
```

0% percentage of the images in human_files_short have a detected dog

94% percentage of the images in dog_files_short have a detected dog

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [11]: import os
         from torchvision import datasets
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         # Variables

         batch_sizes = 32
         num_worker = 0

         transform = transforms.Compose([transforms.Resize(224),
                                         transforms.CenterCrop(224),
                                         transforms.RandomRotation(30),
                                         transforms.RandomHorizontalFlip(),
                                         transforms.ToTensor(),
                                         transforms.Normalize((0.485, 0.456, 0.406),
                                                                (0.229, 0.224, 0.225))
                                         ])

         #Training Image set
         train_set = datasets.ImageFolder('/data/dog_images/train/', transform = transform)

         # Validation Image Set
         vaild_set = datasets.ImageFolder('/data/dog_images/valid', transform = transform)

         # Test Image set
         test_set = datasets.ImageFolder('/data/dog_images/test', transform = transform)

         #Loaders
         #Train_loader
         train_loader = torch.utils.data.DataLoader(train_set, batch_size = batch_sizes, shuffle = True)

         #valid_loader
         valid_loader = torch.utils.data.DataLoader(vaild_set, batch_size= batch_sizes, shuffle = True)

         #Test_loader
         test_loader = torch.utils.data.DataLoader(test_set, batch_size = batch_sizes, shuffle = True)

In [ ]: #Vizulize the data
         #Test for Loader
```

```

images, labels = next(iter(train_loader))
images = images.numpy()
#Plot the figures
fig = plt.figure(figsize=(25, 4))
for i in np.arange(20):
    ax = fig.add_subplot(2, 20/2, i+1, xticks=[], yticks=[])
    # images = images / 2 + 0.5 # unnormalize
    plt.imshow(np.transpose(images[i], (1, 2, 0)))
    #set the title
    # ax.set_title(labels[i])# Classes not defined

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

- How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why?

---> I used to transforms.Resize(224), transforms.CenterCrop(224) to crop & shape the image in the size of 224 x 224.

- Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

---> **Yes!!** By using transforms.RandomRotation(30),transforms.RandomHorizontalFlip(). It helps your model learn better.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [12]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        #Convolutional Layers
        #Sees 224x224x3 Input Image
        self.conv1 = nn.Conv2d(3, 32, 3, stride = 2, padding = 1)
        #Sees 112x112x32
        self.conv2 = nn.Conv2d(32, 64, 3, stride = 2, padding = 1)
        #Sees 56x56x64

```

```

        self.conv3 = nn.Conv2d(64, 128, 3, padding = 1)
        #Sees 28x28x128
#Pooling layers
        self.pool = nn.MaxPool2d(2, 2)
        breeds = 133
#Fully Connected Layers
        self.fc1 = nn.Linear(128 * 7 * 7 , 1024)
        self.fc2 = nn.Linear(1024, breeds )

#Droup probability of 25%
        self.dropout = nn.Dropout(0.25)
def forward(self, x):
    ## Define forward behavior
    x = F.relu(self.conv1(x))
    x = self.pool(x)
    x = F.relu(self.conv2(x))
    x = self.pool(x)

    x = F.relu(self.conv3(x))
    x = self.pool(x)
    #Flatten the Image
    x = x.view(-1, 128 * 7 * 7)

    #To the Fully Connected layer
    x = self.dropout(x)
    x = F.relu(self.fc1(x))
    x = self.dropout(x)
    x = self.fc2(x)

    return x

### You do NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if torch.cuda.is_available():
    model_scratch.cuda()

```

In [13]: print (model_scratch)

```

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

```

```

(fc1): Linear(in_features=6272, out_features=1024, bias=True)
(fc2): Linear(in_features=1024, out_features=133, bias=True)
(dropout): Dropout(p=0.25)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

I decide to use CNN model with three Convolutional layers with Relu activation function and a pooling layer attached to it after every pass.

After passing through the convolutional layer it is flattened and transferred to the fully connected layers. With 'FC1' 128 * 7 * 7 input_layer and 1024 output layers, Relu activation function followed by a dropout layer (p=0.25) and 'FC2' have 1024 input_layer connected to previous layer and 133 output_layer according to the classes.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [14]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.Adam(model_scratch.parameters(), lr = 0.001)

```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [15]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf

        for epoch in range(1, n_epochs+1):
            # initialize variables to monitor training and validation loss
            train_loss = 0.0
            valid_loss = 0.0

            #####
            # train the model #
            #####
            model.train()
            for batch_idx, (data, target) in enumerate(loader[0]):

```

```

        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            optimizer_scratch.zero_grad()
            output = model_scratch(data)
            loss = criterion_scratch(output, target)
            loss.backward()
            optimizer_scratch.step()
            train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loader[1]):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
        out = model_scratch(data)
        vloss = criterion_scratch(out, target)
        valid_loss += ((1 / (batch_idx + 1)) * (vloss.data - valid_loss))
        ## update the average validation loss

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print ('Validation Loss had decreased {:.6f} --> {:.6f}. Saving Current Model')
    torch.save(model_scratch.state_dict(), 'model_scratch.pt')
    valid_loss_min = valid_loss

# return trained model
return model

# train the model
loader = [train_loader, valid_loader]
model_scratch = train(15, loader, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

```

```

Epoch: 1      Training Loss: 4.782468      Validation Loss: 4.614245
Validation Loss had decreased inf --> 4.614245. Saving Current Model..
Epoch: 2      Training Loss: 4.415905      Validation Loss: 4.315432
Validation Loss had decreased 4.614245 --> 4.315432. Saving Current Model..
Epoch: 3      Training Loss: 4.159312      Validation Loss: 4.141236
Validation Loss had decreased 4.315432 --> 4.141236. Saving Current Model..
Epoch: 4      Training Loss: 3.967436      Validation Loss: 3.981108
Validation Loss had decreased 4.141236 --> 3.981108. Saving Current Model..
Epoch: 5      Training Loss: 3.788776      Validation Loss: 3.893961
Validation Loss had decreased 3.981108 --> 3.893961. Saving Current Model..
Epoch: 6      Training Loss: 3.631158      Validation Loss: 3.883883
Validation Loss had decreased 3.893961 --> 3.883883. Saving Current Model..
Epoch: 7      Training Loss: 3.494131      Validation Loss: 3.937231
Epoch: 8      Training Loss: 3.359942      Validation Loss: 3.780471
Validation Loss had decreased 3.883883 --> 3.780471. Saving Current Model..
Epoch: 9      Training Loss: 3.208518      Validation Loss: 3.774253
Validation Loss had decreased 3.780471 --> 3.774253. Saving Current Model..
Epoch: 10     Training Loss: 3.066980      Validation Loss: 3.747632
Validation Loss had decreased 3.774253 --> 3.747632. Saving Current Model..
Epoch: 11     Training Loss: 2.950228      Validation Loss: 3.829823
Epoch: 12     Training Loss: 2.808523      Validation Loss: 3.814773
Epoch: 13     Training Loss: 2.683963      Validation Loss: 3.845926
Epoch: 14     Training Loss: 2.562822      Validation Loss: 3.841945
Epoch: 15     Training Loss: 2.459389      Validation Loss: 3.930493

```

```

In [16]: # load the model that got the best validation accuracy
         model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```

In [17]: def test(loaders, model, criterion, use_cuda):

         # monitor test loss and accuracy
         test_loss = 0.
         correct = 0.
         total = 0.

         model.eval()
         for batch_idx, (data, target) in enumerate(loaders[0]):
             # move to GPU
             if use_cuda:
                 data, target = data.cuda(), target.cuda()
             # forward pass: compute predicted outputs by passing inputs to the model
             output = model(data)

```



```

        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    loader = [test_loader]
    test(loader, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.719241

Test Accuracy: 14% (121/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [18]: ## TODO: Specify data loaders
        batch_sizes = 32
        num_worker = 0

        transform = transforms.Compose([transforms.Resize(224),
                                        transforms.CenterCrop(224),
                                        transforms.RandomRotation(30),

                                        transforms.RandomHorizontalFlip(p=0.2),
                                        transforms.ToTensor(),

```

```

        transforms.Normalize((0.485, 0.456, 0.406),
                              (0.229, 0.224, 0.225))
    ])

#Training Image set
train_set_tran = datasets.ImageFolder('/data/dog_images/train/', transform = transform)

# Validation Image Set
vaild_set_tran = datasets.ImageFolder('/data/dog_images/valid', transform = transform)

# Test Image set
test_set_tran = datasets.ImageFolder('/data/dog_images/test', transform = transform)

#Loaders
#Train_loader
train_loader_tran = torch.utils.data.DataLoader(train_set_tran, batch_size = batch_size)

#valid_loader
valid_loader_tran = torch.utils.data.DataLoader(vaild_set_tran, batch_size= batch_sizes)

#Test_loader
test_loader_tran = torch.utils.data.DataLoader(test_set_tran, batch_size = batch_sizes,

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [19]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture
model_transfer = models.resnet50(pretrained=True)
#print (ResNet_transfer)
#Hyperparamters

input_layer = model_transfer.fc.in_features
output_layer = 133

#Freeze the parameters in Resnet!! Okay!!

for prams in model_transfer.parameters():
    prams.requires_grad = False

#define the custom layer for resnet
fc_custom = nn.Linear(input_layer, output_layer)

```

```
#Apply the custom layer
```

```
model_transfer.fc = fc_custom
```

```
#move model to GPU!!
```

```
if use_cuda:
```

```
    model_transfer = model_transfer.cuda()
```

Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/100%|| 102502400/102502400 [00:01<00:00, 71796408.28it/s]

In [20]:

```
#Unfreeze the custom layer to adapt weights!!
```

```
for prams in model_transfer.fc.parameters():
```

```
    prams.requires_grad = True
```

```
print (model_transfer)
```

ResNet(

(conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)

(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(relu): ReLU(inplace)

(maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)

(layer1): Sequential(

(0): Bottleneck(

(conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(relu): ReLU(inplace)

(downsample): Sequential(

(0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

)

)

(1): Bottleneck(

(conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

(relu): ReLU(inplace)

```

)
(2): Bottleneck(
  (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
(layer2): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (2): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (3): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
    )
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
  (1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
    )
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (2): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)

```

```

    )
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=2048, out_features=133, bias=True)
)

```

```
In [21]: print (model_transfer.fc)
```

```
Linear(in_features=2048, out_features=133, bias=True)
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

It used ResNet50 as my model and replaced the fc layer with my custom made layer suited for the given task. Freezed all the parameters and unfreezed the custom parameters.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [22]: criterion_transfer = nn.CrossEntropyLoss()

ResNetParams = model_transfer.fc.parameters()

optimizer_transfer = optim.Adam(ResNetParams, lr = 0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [23]: # train the model
loaders = [train_loader_tran, valid_loader_tran]

def train_transfer(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
```

```

#####
model.train()
for batch_idx, (data, target) in enumerate(loaders[0]):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
        ## find the loss and update the model parameters accordingly
        optimizer_transfer.zero_grad()
        output = model_transfer(data)
        loss = criterion_transfer(output, target)
        loss.backward()
        optimizer_transfer.step()
        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
        ## record the average training loss, using something like
        ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders[1]):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
        out = model_transfer(data)
        vloss = criterion_transfer(out, target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (vloss.data - valid_loss))
        ## update the average validation loss

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print ('Validation Loss had decreased {:.6f} --> {:.6f}. Saving Current Model')
    torch.save(model_transfer.state_dict(), 'model_transfer.pt')
    valid_loss_min = valid_loss

# return trained model
return model

```



```

model_transfer = train_transfer(25, loaders, model_transfer, optimizer_transfer, crite

# load the model that got the best validation accuracy (uncomment the line below)
#model_transfer.load_state_dict(torch.load('model_transfer.pt'))

```

```

Epoch: 1      Training Loss: 2.328061      Validation Loss: 1.025582
Validation Loss had decreased inf --> 1.025582. Saving Current Model..
Epoch: 2      Training Loss: 0.879667      Validation Loss: 0.761514
Validation Loss had decreased 1.025582 --> 0.761514. Saving Current Model..
Epoch: 3      Training Loss: 0.642100      Validation Loss: 0.655831
Validation Loss had decreased 0.761514 --> 0.655831. Saving Current Model..
Epoch: 4      Training Loss: 0.561082      Validation Loss: 0.660429
Epoch: 5      Training Loss: 0.490589      Validation Loss: 0.633445
Validation Loss had decreased 0.655831 --> 0.633445. Saving Current Model..
Epoch: 6      Training Loss: 0.443461      Validation Loss: 0.597901
Validation Loss had decreased 0.633445 --> 0.597901. Saving Current Model..
Epoch: 7      Training Loss: 0.423007      Validation Loss: 0.604735
Epoch: 8      Training Loss: 0.377476      Validation Loss: 0.620995
Epoch: 9      Training Loss: 0.360276      Validation Loss: 0.602939
Epoch: 10     Training Loss: 0.335009      Validation Loss: 0.614755
Epoch: 11     Training Loss: 0.338916      Validation Loss: 0.583603
Validation Loss had decreased 0.597901 --> 0.583603. Saving Current Model..
Epoch: 12     Training Loss: 0.323460      Validation Loss: 0.631725
Epoch: 13     Training Loss: 0.291297      Validation Loss: 0.617680
Epoch: 14     Training Loss: 0.295719      Validation Loss: 0.771904
Epoch: 15     Training Loss: 0.282835      Validation Loss: 0.684667
Epoch: 16     Training Loss: 0.261980      Validation Loss: 0.643550
Epoch: 17     Training Loss: 0.268791      Validation Loss: 0.644148
Epoch: 18     Training Loss: 0.256880      Validation Loss: 0.640594
Epoch: 19     Training Loss: 0.251191      Validation Loss: 0.667173
Epoch: 20     Training Loss: 0.235856      Validation Loss: 0.648816
Epoch: 21     Training Loss: 0.241209      Validation Loss: 0.598467
Epoch: 22     Training Loss: 0.237662      Validation Loss: 0.662421
Epoch: 23     Training Loss: 0.224308      Validation Loss: 0.702431
Epoch: 24     Training Loss: 0.213824      Validation Loss: 0.704913
Epoch: 25     Training Loss: 0.211493      Validation Loss: 0.671364

```

```

In [24]: # load the model that got the best validation accuracy
         model_transfer.load_state_dict(torch.load('model_trasfer.pt'))

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [25]: loaders_transfer = [train_loader_tran]
        test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.204092

Test Accuracy: 93% (6248/6680)

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [26]: ### TODO: Write a function that takes a path to an image as input
        ### and returns the dog breed that is predicted by the model.

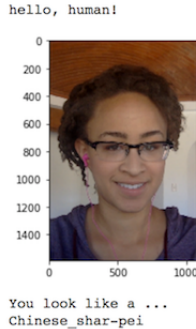
        # list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in train_set_tran.classes]
# Read Images
def image_read(img_path):
    image = Image.open(img_path)
    image = transform(image)
    image = image.unsqueeze(0)
    if use_cuda:
        image = image.cuda()
    return image

#print (class_names)
def predict_breed_transfer(img_path):
    # load the image and return the predicted breed

    model_transfer.eval()
    out = model_transfer(image_read(img_path))
    preds = torch.argmax(out)
    return class_names[preds]

In [27]: print (predict_breed_transfer('/data/dog_images/train/001.Affenpinscher/Affenpinscher_0
        print (predict_breed_transfer('/data/dog_images/train/039.Bull_terrier/Bull_terrier_027
        print (predict_breed_transfer('/data/dog_images/valid/113.Old_english_sheepdog/Old_engl
        print (predict_breed_transfer('/data/dog_images/test/068.Flat-coated_retriever/Flat-coa
```

```
Affenpinscher
Bull terrier
Old english sheepdog
Flat-coated retriever
```



Sample Human Output

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

In [32]: *### TODO: Write your algorithm.*

Feel free to use as many code cells as needed.

```
def run_app(img_path):


    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()

    if dog_detector(img_path) == True:
        pred = predict_breed_transfer (img_path)
        print ('It is a dog. \nIts predicted breed is... {}'.format (pred))
    elif face_detector(img_path) > 0:
        pred = predict_breed_transfer (img_path)
        print ('Hi, Human \nYou look like... \n{}'.format(pred))
    else:
        print ('Sorry!! I dont know this.')
    ## handle cases for a human face, dog, and neither
```

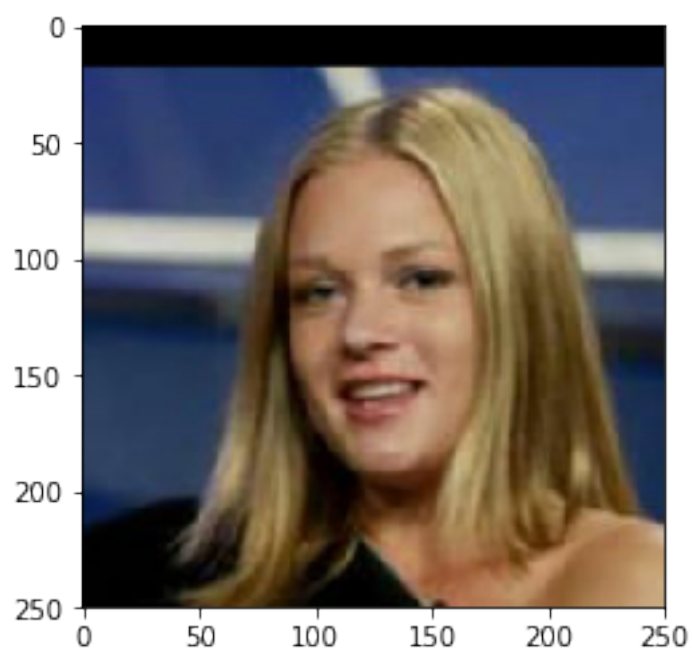
In [33]: `print (run_app('/data/lfw/Abba_Eban/Abba_Eban_0001.jpg'))`

`print (run_app('/data/lfw/AJ_Cook/AJ_Cook_0001.jpg'))`


`print (run_app('/data/dog_images/test/009.American_water_spaniel/American_water_spaniel'))`



Hi, Human
You look like...
American water spaniel
None



Hi, Human
You look like...
Silky terrier
None



It is a dog.
Its predicted breed is... American water spaniel
None

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!


Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

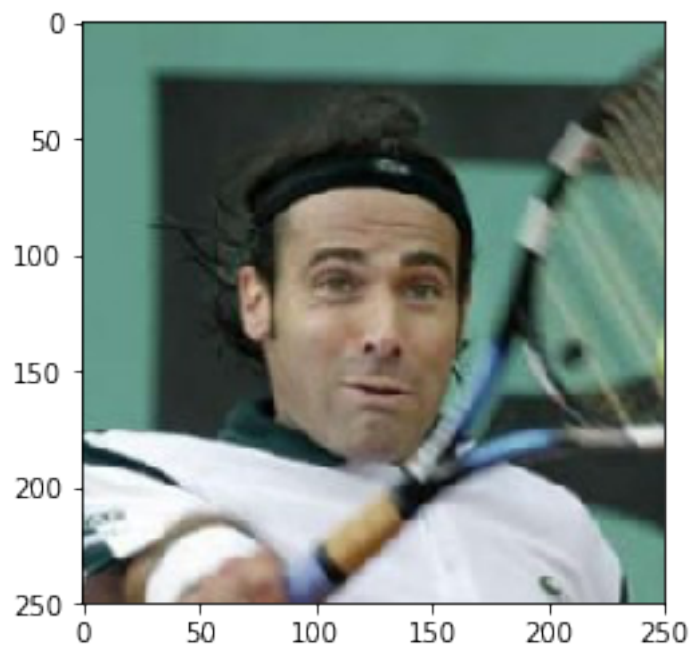
Answer: (Three possible points for improvement)

- 1) Can augment the data, and more data can definately help improve the accuracy.
- 2) Choosing different hyperparameters: different optimizer, dropout, weight initialization.
- 3) Choosing a different model can all surely help in improving model. :)


```
In [34]: ## TODO: Execute your algorithm from Step 6 on  
## at least 6 images on your computer.  
## Feel free to use as many code cells as needed.  
  
## suggested code, below  
for file in np.hstack((human_files[:3], dog_files[:3])):  
    run_app(file)
```




```
Hi, Human  
You look like...  
American staffordshire terrier
```



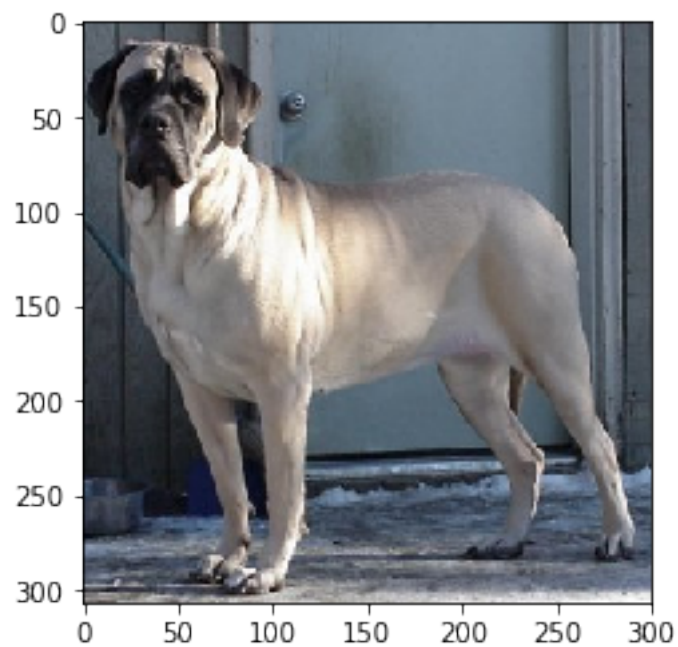
Hi, Human
You look like...
Australian shepherd




Hi, Human
You look like...
American water spaniel



It is a dog.
Its predicted breed is... Mastiff




It is a dog.
Its predicted breed is... Bullmastiff




It is a dog.
Its predicted breed is... Bullmastiff

```
In [35]: run_app ('images/Curly-coated_retriever_03896.jpg')
```



It is a dog.
Its predicted breed is... Curly-coated retriever

```
In [37]: run_app('images/Brittany_02625.jpg')
```



It is a dog.
Its predicted breed is... Brittany

In []: